

---

# **Lab 3**

**Blake Smith**

**May 30, 2020**

CONTENTS

1 Sequential Map Reduce, A Baseline 1

1.1 The Phases of Map Reduce (In the Sequential Case) 1

1.2 Modifications to *mrsequential.go* 1

## **SEQUENTIAL MAP REDUCE, A BASELINE**

Before proceeding with testing my implementation from Lab 2 I will first gather some metrics from the sequential implementation given in *mrsequential.go*. This baseline will be used to determine at what point the multi-process implementation becomes worthwhile and what the trade offs are in regards to space & time complexity.

### **1.1 The Phases of Map Reduce (In the Sequential Case)**

In order to estimate the performance of a sequential map-reduce I will be collecting the following information at each step in the execution.

1. Read input files and pass into the map function, producing a collection of intermediate values.
  - Time taken to read in the input files and produce the intermediate collection.
  - Space required to store the intermediate values
  - Time taken to sort the intermediate values
2. Group the intermediate values by key, producing a list of values for every key.
  - Time taken to group the values by key
  - Amount of memory used in that process
3. Run Reduce on each key and create a single output file
  - Time taken to complete all reduce jobs and produce the full output

### **1.2 Modifications to *mrsequential.go***